

# Challenges of Frequent Pattern Mining

---

- Challenges
  - Multiple scans of a transaction database
  - Huge number of candidates
  - Tedious workload of support counting for candidates
- Improving Apriori: general ideas
  - Reduce the number of transaction database scans
  - Shrink the number of candidates
  - Facilitate support counting of candidates




# Partition: Scan Database Only Twice

---

- Approach

- Divide a database into k pieces (local databases called *partition*)



- Each partition should reside in main memory
- Find *local frequent patterns* in each partition (scan 1)
  - localMinSup is set as  $(\text{minSup} / k)$
  - Local frequent patterns have their localSup larger than localMinSup in any local database
- Consolidate global frequent patterns (scan 2)



# Partition: Scan Database Only Twice

---

- Guarantee that frequent patterns are never missed
    - Any itemset potentially frequent in DB must be frequent in *at least one partition* of DB
- |    |    |    |    |
|----|----|----|----|
| P1 | P2 | P3 | P4 |
|----|----|----|----|
- localMinSup is set as  $(\text{minSup} / k)$
  - Local frequent patterns have their localSup larger than localMinSup in any local database
- A. Savasere, E. Omiecinski, and S. Navathe. *An efficient algorithm for mining association in large databases*. In *VLDB'95*



# DHP: Reduce the Number of Candidates

- Use a hash table for  $(k+1)$ -itemsets during determining  $k$ -itemsets by database scan
  - Candidates of 1-itemset: a, b, c, d, e, f, ....
    - What if 10,000 items? => 100,000,000 candidate 2-itemsets!
  - Hash table for 2-itemsets: {ab, ad, ae} {bd, be, de}  
...
  - A  $(k+1)$ -itemset whose corresponding **hash bucket count** is below the threshold cannot be frequent
  - ab is not a candidate 2-itemset if the sum of count of {ab, ad, ae} is below threshold of minimum support (say, 50)
    - Effective in reducing # of candidate frequent 2-itemsets
- J. Park, M. Chen, and P. Yu. An effective hash-based algorithm for mining association rules. In *SIGMOD'95*

count	itemsets
35	{ab, ad, ae}
88	{bd, be, de}
.	.
.	.
.	.
102	{yz, qs, wt}

**Hash Table**



# Sampling for Frequent Patterns

- Select a sample of an original database, mine frequent patterns within sample using Apriori (in the same way as before)



Sampling =>



sampled DB (SDB)

- Use a smaller value of the minimum support for a sample (say,  $\text{minSup}/4$ )
- Problems with the simple sampling
  - Some of frequent patterns found in SDB (i.e., **S**) are not really frequent in the original database
  - Some of true frequent patterns could be missed if they are not included in **S**

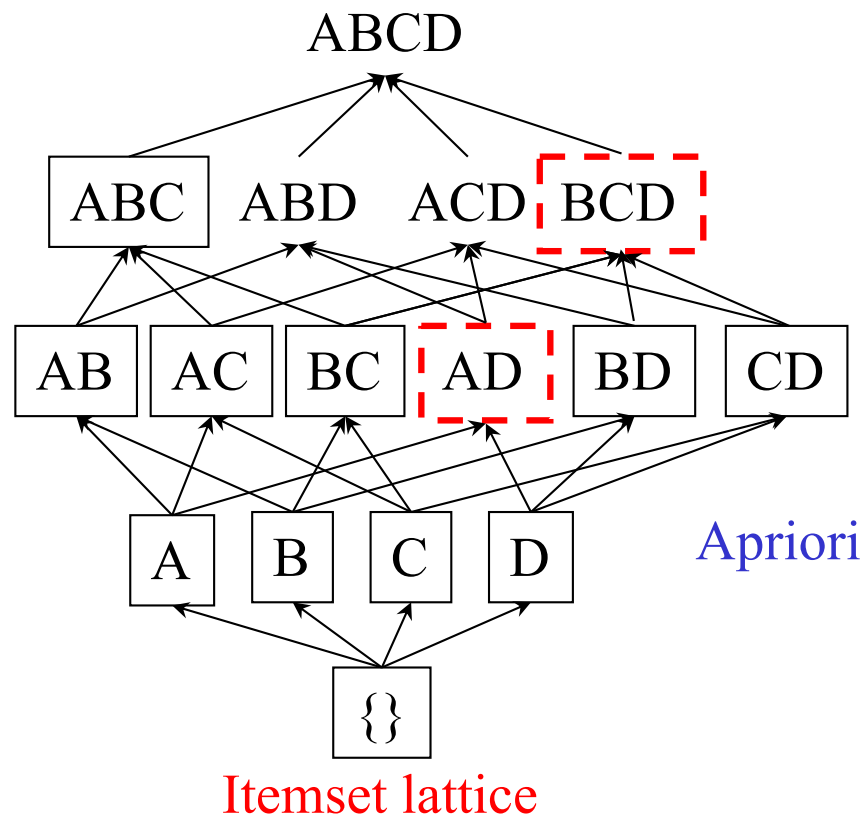


# Sampling for Frequent Patterns

- Solutions: two more scanning for verification
- Scan the whole database once
  - Verify a collection of frequent itemsets,  $S$ , found in sample, and its negative borders (NB: not in  $S$ , but all its subsets in  $S$ )
    - $S = \{a\}, \{b\}, \{c\}, \{f\}, \{a,b\}, \{a,c\}, \{a,f\}, \{c,f\}, \{a,c,f\}$
    - NB =  $\{b,c\}, \{b,f\}, \{d\}, \{e\}$
- Scan the whole database again
  - Find missed frequent patterns (due to the success of NBs)
- H. Toivonen. Sampling large databases for association rules. In *VLDB'96*



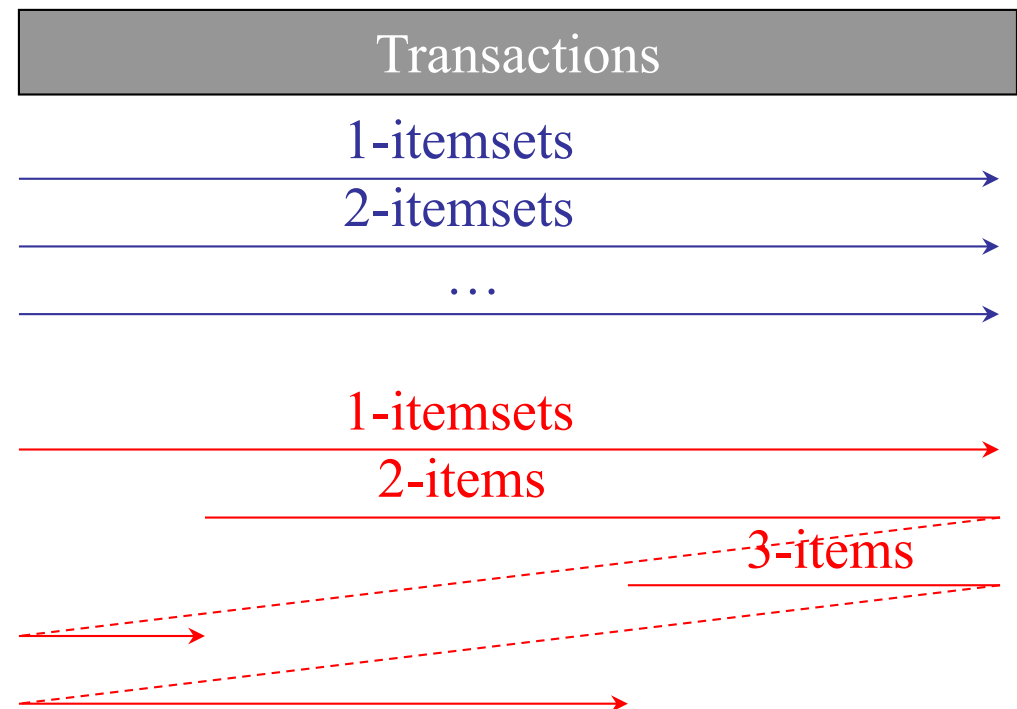
# DIC: Reduce Number of Scans



S. Brin R. Motwani, J. Ullman,  
and S. Tsur. *Dynamic itemset  
counting and implication rules for  
market basket data*. In  
*SIGMOD'97*

March 24, 2020

- Once both A and D are determined frequent, the counting of AD begins
- Once all length-2 subsets of BCD are determined frequent, the counting of BCD begins



DIC



# Bottleneck of Frequent-pattern Mining

---

- Multiple database scans are **costly**
- Mining long patterns needs many passes of scanning and generates lots of candidates
  - To find frequent itemset  $i_1 i_2 \dots i_{100}$ 
    - # of scans: **100**
    - # of Candidates:  $\binom{100}{1} + \binom{100}{2} + \dots + \binom{100}{100} = 2^{100} - 1 = \mathbf{1.27 * 10^{30} !}$
- Bottleneck: candidate-generation-and-test
- Can we avoid candidate generation?





# FP-Growth: Mining Frequent Patterns Without Candidate Generation

---

- Grow long patterns from short ones using local frequent items
  - “abc” is a frequent pattern
  - Get all transactions having “abc”
    - Denoted as  $DB|abc$
  - “d” is a local frequent item in  $DB|abc \rightarrow abcd$  is a frequent pattern



# FP-Growth: Construct FP-tree from a Transaction Database

<i>TID</i>	<i>Items bought</i>	<i>(ordered) frequent items</i>
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o, w}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}


*min\_support* = 3

1. Scan DB once, find frequent 1-itemset (single item pattern)
2. Sort frequent items in frequency descending order, f-list
3. Scan DB again, construct FP-tree

**Header Table**

<i>Item</i>	<i>frequency</i>	<i>head</i>
f	4	
c	4	
a	3	
b	3	
m	3	
p	3	

**F-list**=f-c-a-b-m-p



# Benefits of the FP-tree Structure

---

- Completeness
  - Preserve *complete (i.e., lossless) information* for frequent pattern mining
  - Never break a long pattern of any transaction
- Compactness
  - Remove irrelevant info—infrequent items are gone
  - Items in frequency descending order: the more frequently occurring, the more likely to be shared
  - Never be larger than the original database
    - (not counting node-links and the *count* field)
    - For Connect-4 DB, compression ratio could be over 100



# Partition Patterns and Databases

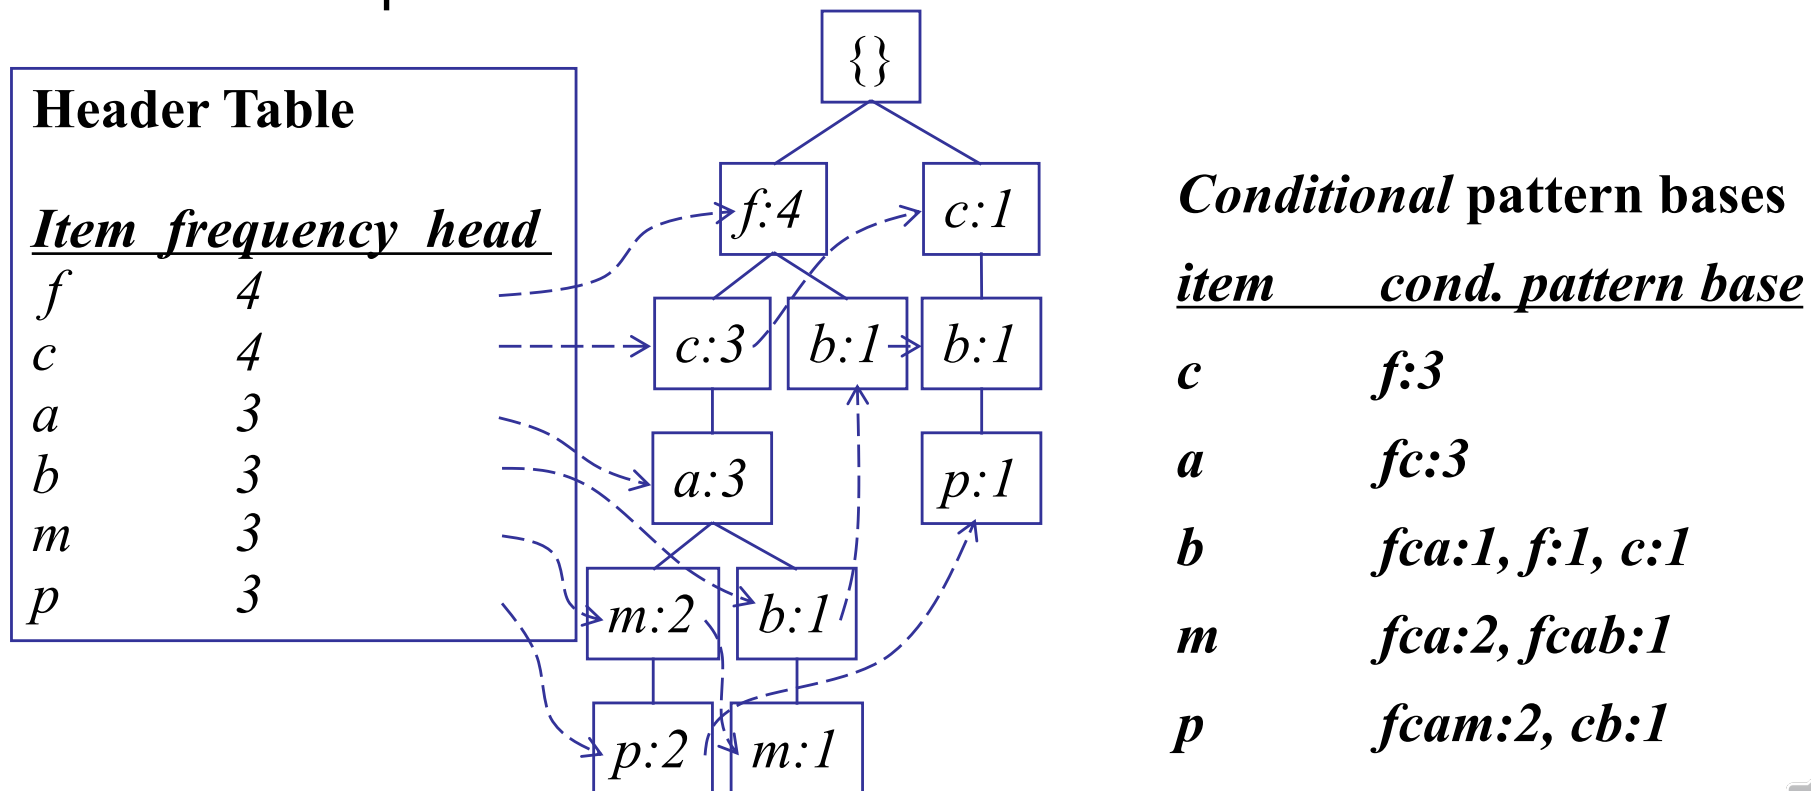
---

- Frequent patterns can be *partitioned* into (*disjoint*) subsets according to f-list
  - F-list=f-c-a-b-m-p
  - Patterns containing p
  - Patterns having m but no p
  - Patterns having m but no m nor p (i.e., not containing m and p)
  - ...
  - Patterns having c but no a nor b, m, p
  - Pattern f
- Completeness and non-redundancy




# Find Patterns Having P From P-conditional Database

- Starting at the frequent item header table in the FP-tree
- Traverse the FP-tree by following the link of each frequent item  $p$
- Accumulate all of *transformed prefix paths* of item  $p$  to form  $p$ 's conditional pattern base

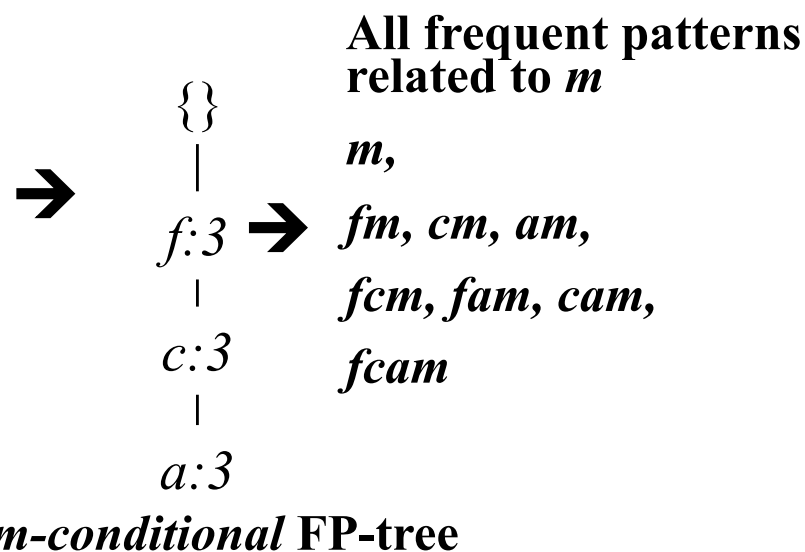


# From Conditional Pattern-bases to Conditional FP-trees

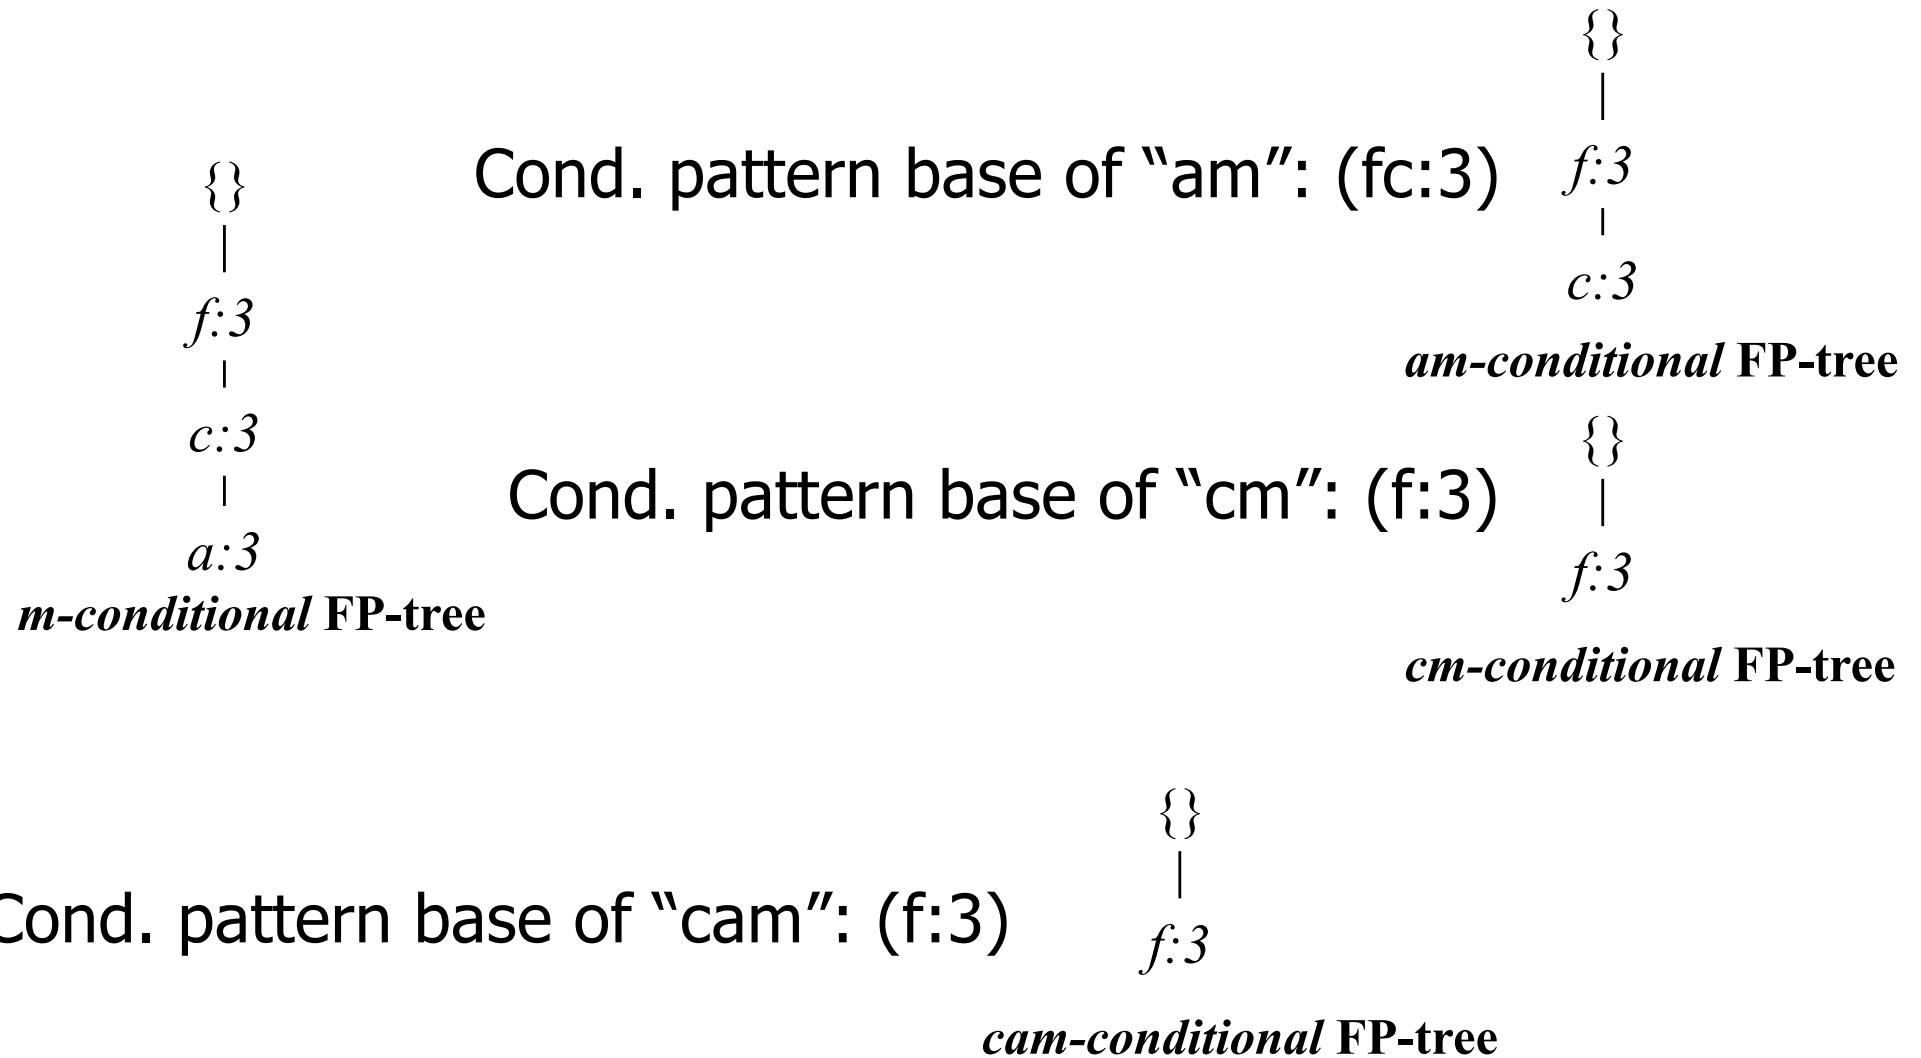
- For each conditional database (i.e., pattern-base)
    - Accumulate the count for each item in the base
    - Construct the FP-tree for the frequent items of the pattern base
- min\_support = 3*



*m*-conditional pattern base:  
*fca:2, fcab:1*

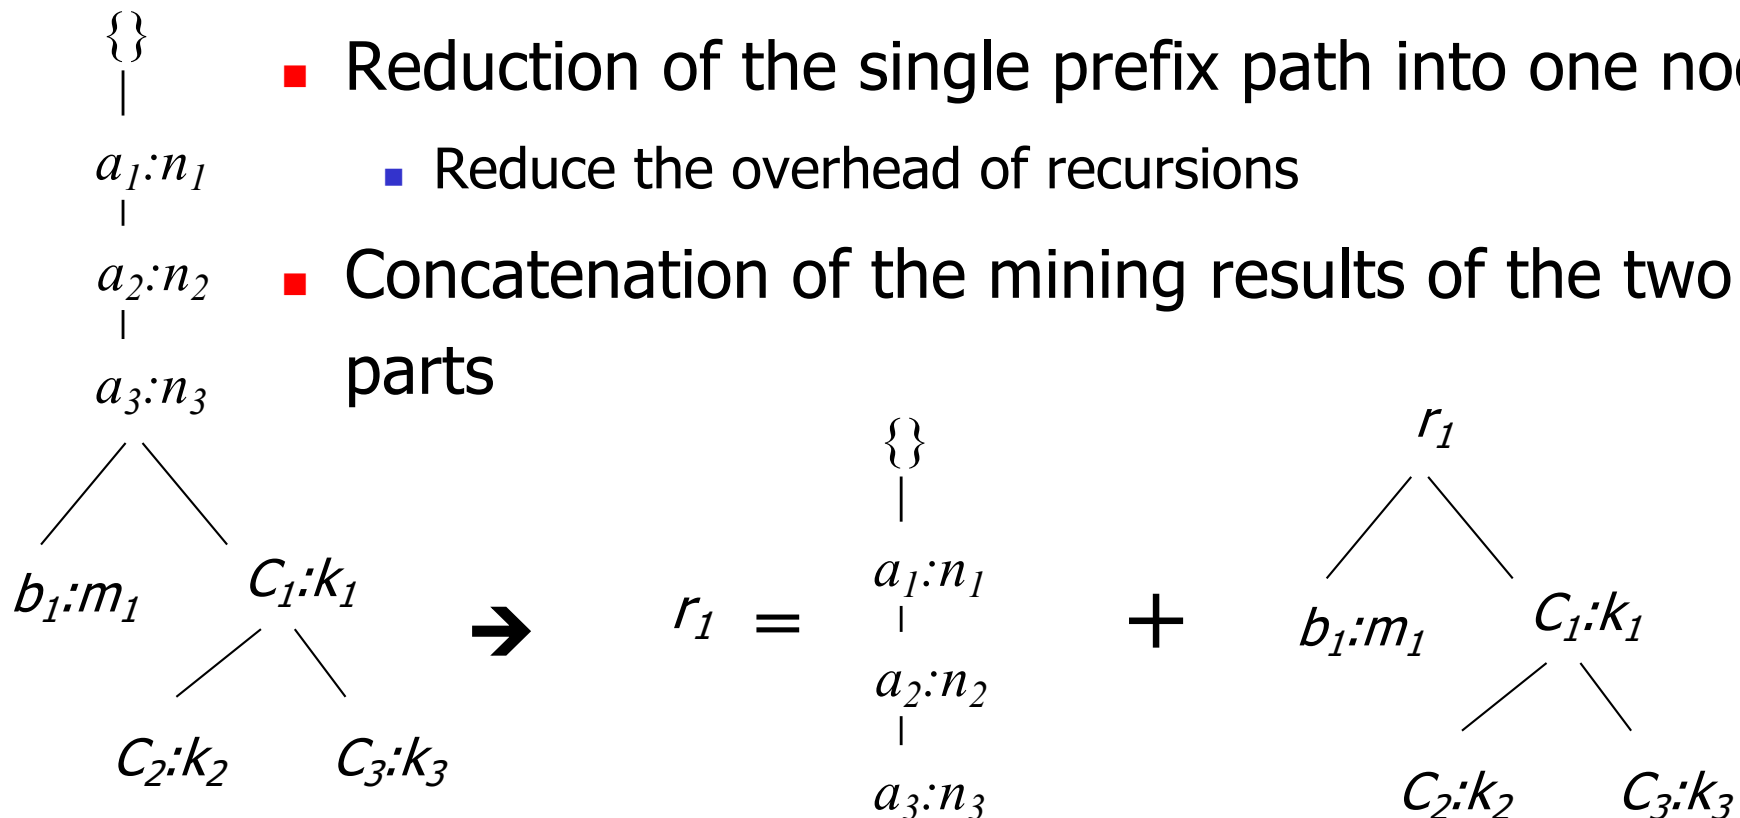


# Recursion: Mining Each Conditional FP-tree



# A Special Case: Single Prefix Path in FP-tree

- Suppose a (conditional) FP-tree T has a shared single prefix-path P
- Mining can be decomposed into two parts
  - Reduction of the single prefix path into one node
    - Reduce the overhead of recursions
  - Concatenation of the mining results of the two parts






# Summary of Ideas with FP-Growth

---

- Idea: Frequent pattern growth
  - Grow frequent patterns by adding a new frequent item recursively
- Method
  - For each frequent item, construct its conditional pattern-base, and then its conditional FP-tree
  - Repeat the process on each newly created conditional FP-tree
  - Until the resulting FP-tree is empty, or it contains only a single path
    - A single path will generate all the combinations of its sub-paths
    - Each of the combinations is a frequent pattern



# FP-Growth vs. Apriori: Scalability With the Support Threshold



# Why Is FP-Growth the Winner?

---

- Divide-and-conquer:
  - decompose both the mining task and a database according to the frequent patterns obtained so far
  - leads to focused search of smaller databases
- Other factors
  - no candidate generation and no candidate test
  - compressed database: FP-tree structure
  - no repeated scans of the entire database: just twice
  - basic operations
    - counting local frequent items and building a sub FP-tree
    - no pattern search and matching

