Chapter 5: Mining Frequent Patterns, Association and Correlations

- Basic concepts and a road map
- Efficient and scalable frequent itemset mining methods
- Mining various kinds of association rules
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

Constraint-based (Query-Directed) Mining

- Finding all the patterns in a database autonomously? — unrealistic!
  - The patterns could be too many but not focused!
- Data mining should be an interactive process
  - User directs what to be mined
  - She/he uses a data mining query language (or a graphical user interface)
- Constraint-based mining
  - User flexibility: provides constraints on what to be mined
  - System optimization: explores such constraints for efficient mining—constraint-based mining

Constraints in Data Mining

- Knowledge type constraint:
  - classification, association, clustering etc.
- Data constraint — using SQL-like queries
  - find product pairs sold together in stores in Chicago in Dec.'18
- Dimension/level constraint
  - in relevance to region, price, brand, customer category
- Interestingness constraint
  - strong rules: min_support $\geq$ 3%, min_confidence $\geq$ 60%

Constrained Mining vs. Other Operations

- Constrained mining vs. constraint-based search
  - Both are aimed at reducing search space
    - Constrained mining: finding all patterns satisfying constraints
    - Constraint-based search: finding some (or one) answer in constraint-based search in AI

Constrained Mining vs. Other Operations

- Constrained mining vs. query processing in DBMS
  - Both are aimed at finding all answers
    - Query processing: finding tuples in a database
    - Constrained mining: discovering patterns hidden in a database
  - Constrained mining shares a similar philosophy as pushing selections deeply in query processing

Anti-Monotonicity in Constraint Pushing

- Anti-monotonicity on a constraint
  - *When an itemset S **violates** the constraint, so does any of its superset*
  - *sum(S.Price)* $\leq v$ is anti-monotone
  - *sum(S.Price)* $\geq v$ is not anti-monotone
- Example. C: range(S.profit) $\leq$ 15 is anti-monotone
  - Itemset *ab* violates C
  - So does every superset of *ab*

TDB (min_sup=2)

TID	Transaction
10	a, b, c, d, f
20	b, c, d, f, g, h
30	a, c, d, e, f
40	c, e, f, g

Item	Profit
a	40
b	0
c	-20
d	10
e	-30
f	30
g	20
h	-10

Monotonicity for Constraint Pushing

TDB (min_sup=2)

TID	Transaction
10	a, b, c, d, f
20	b, c, d, f, g, h
30	a, c, d, e, f
40	c, e, f, g

- Monotonicity on a constraint
  - *When an intemset S **satisfies** the constraint, so does any of its superset*
  - *sum(S.Price)* $\geq v$ is monotone
  - *min(S.Price)* $\leq v$ is monotone
- Example. C: range(S.profit) $\geq$ 15
  - Itemset *ab* satisfies C
  - So does every superset of *ab*

Item	Profit
a	40
b	0
c	-20
d	10
e	-30
f	30
g	20
h	-10

Succinctness

- Succinctness on a constraint
  - When a set of items (say A_1) satisfies a constraint C, any set S satisfying C is 'simply computed' based on A_1 (In this case, S contains a subset belonging to A_1)
    - $min(S.Price) \leq v$ is succinct
    - $sum(S.Price) \geq v$ is not succinct
- Good thing
  - Without looking at the transaction database, whether an itemset S satisfies constraint C can be determined based on the selection of items
- Optimization: If C is succinct, C is pre-counting pushable

The Apriori Algorithm — Example

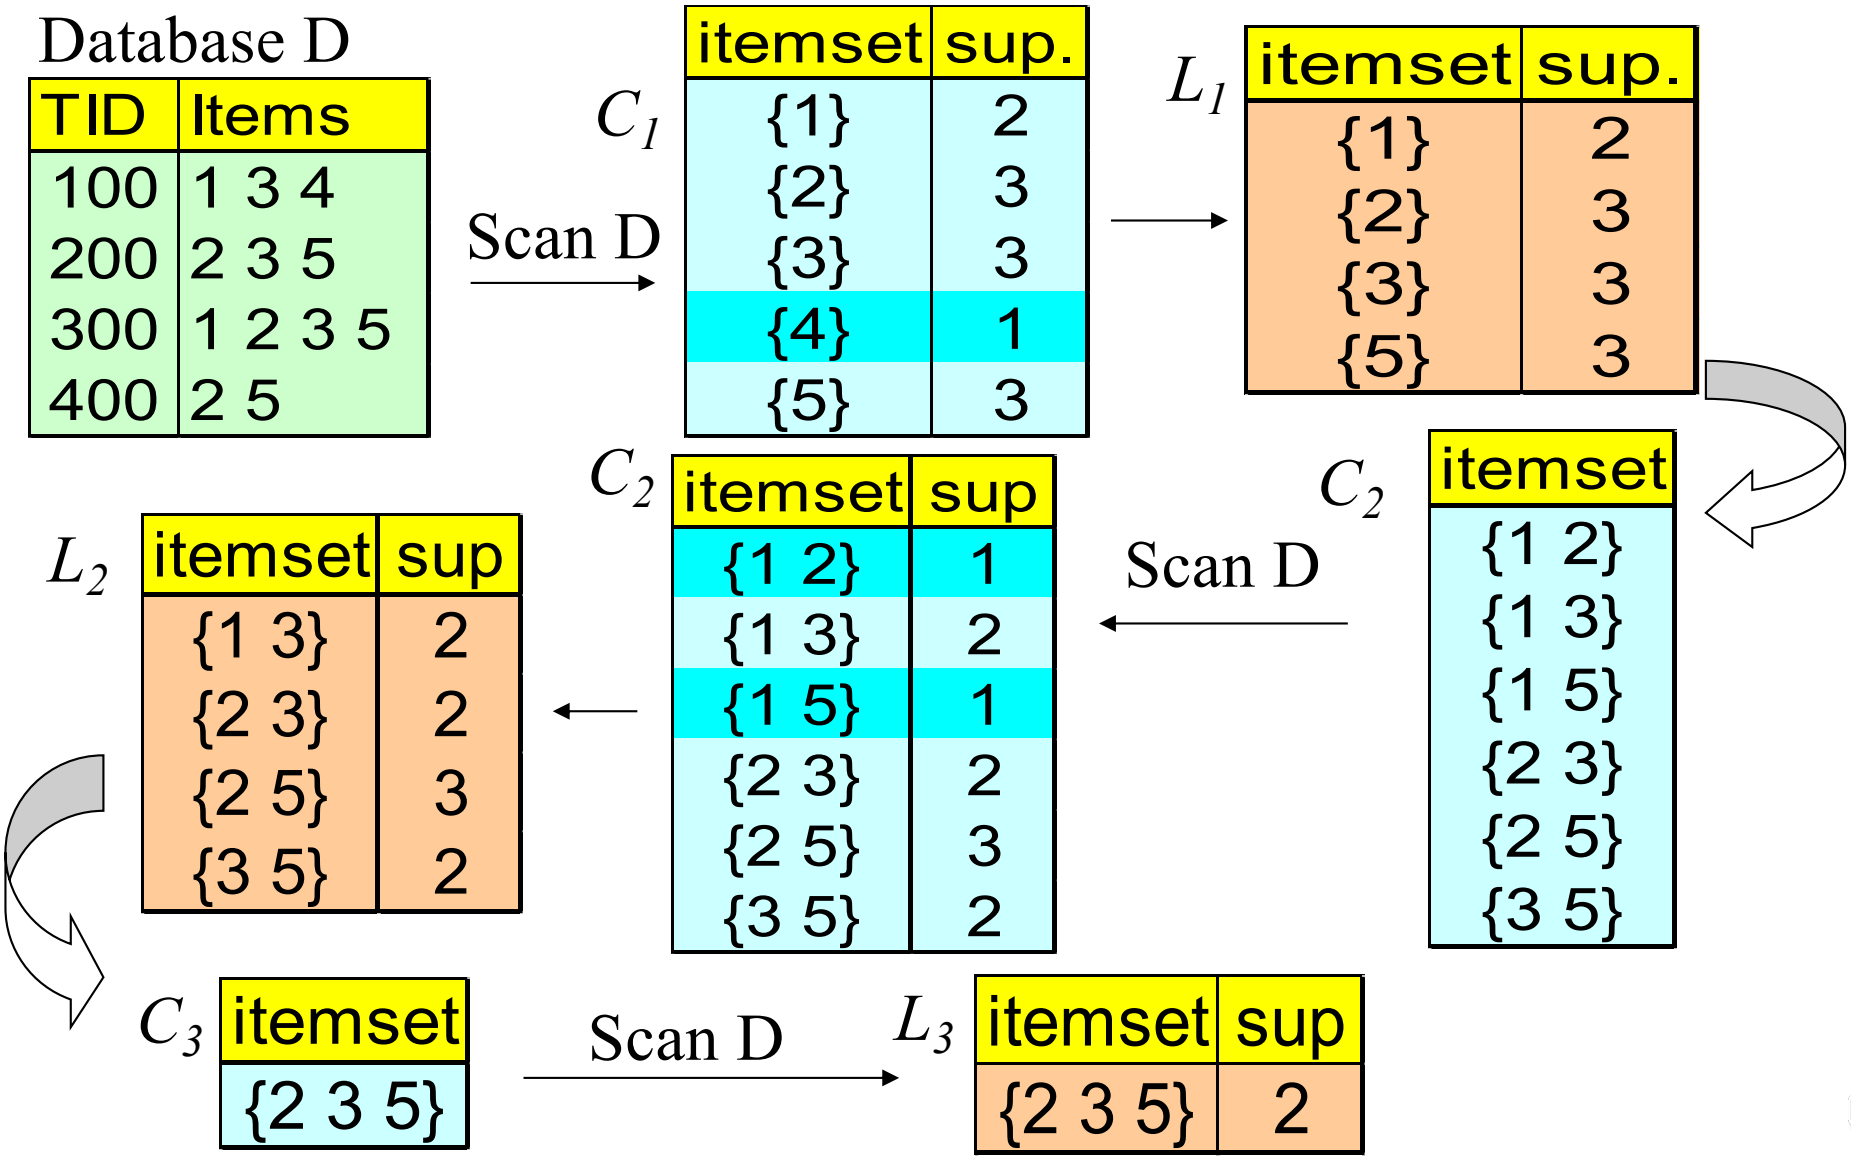

Naïve Algorithm: Apriori + Constraint

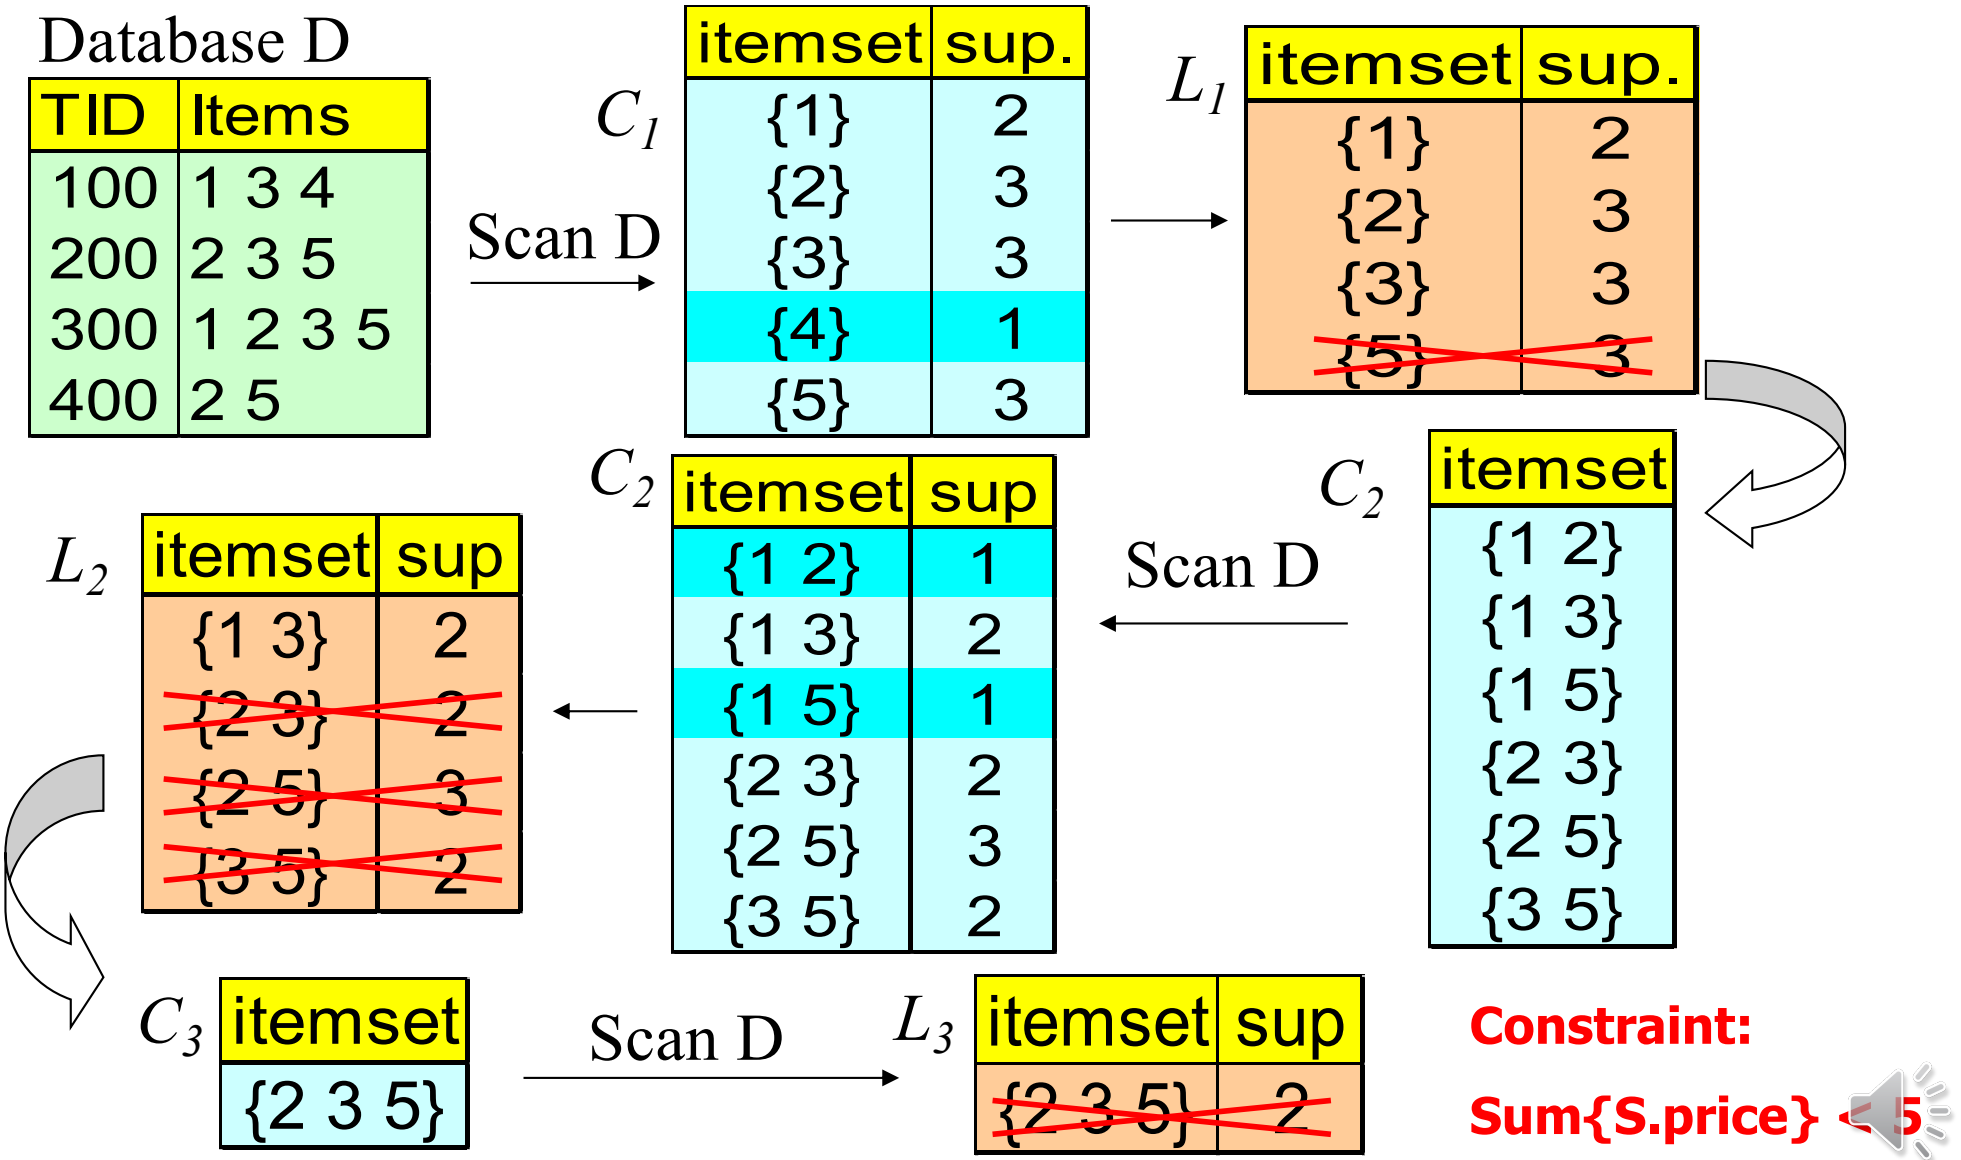

The Constrained Apriori Algorithm: Push an Anti-monotone Constraint Deep

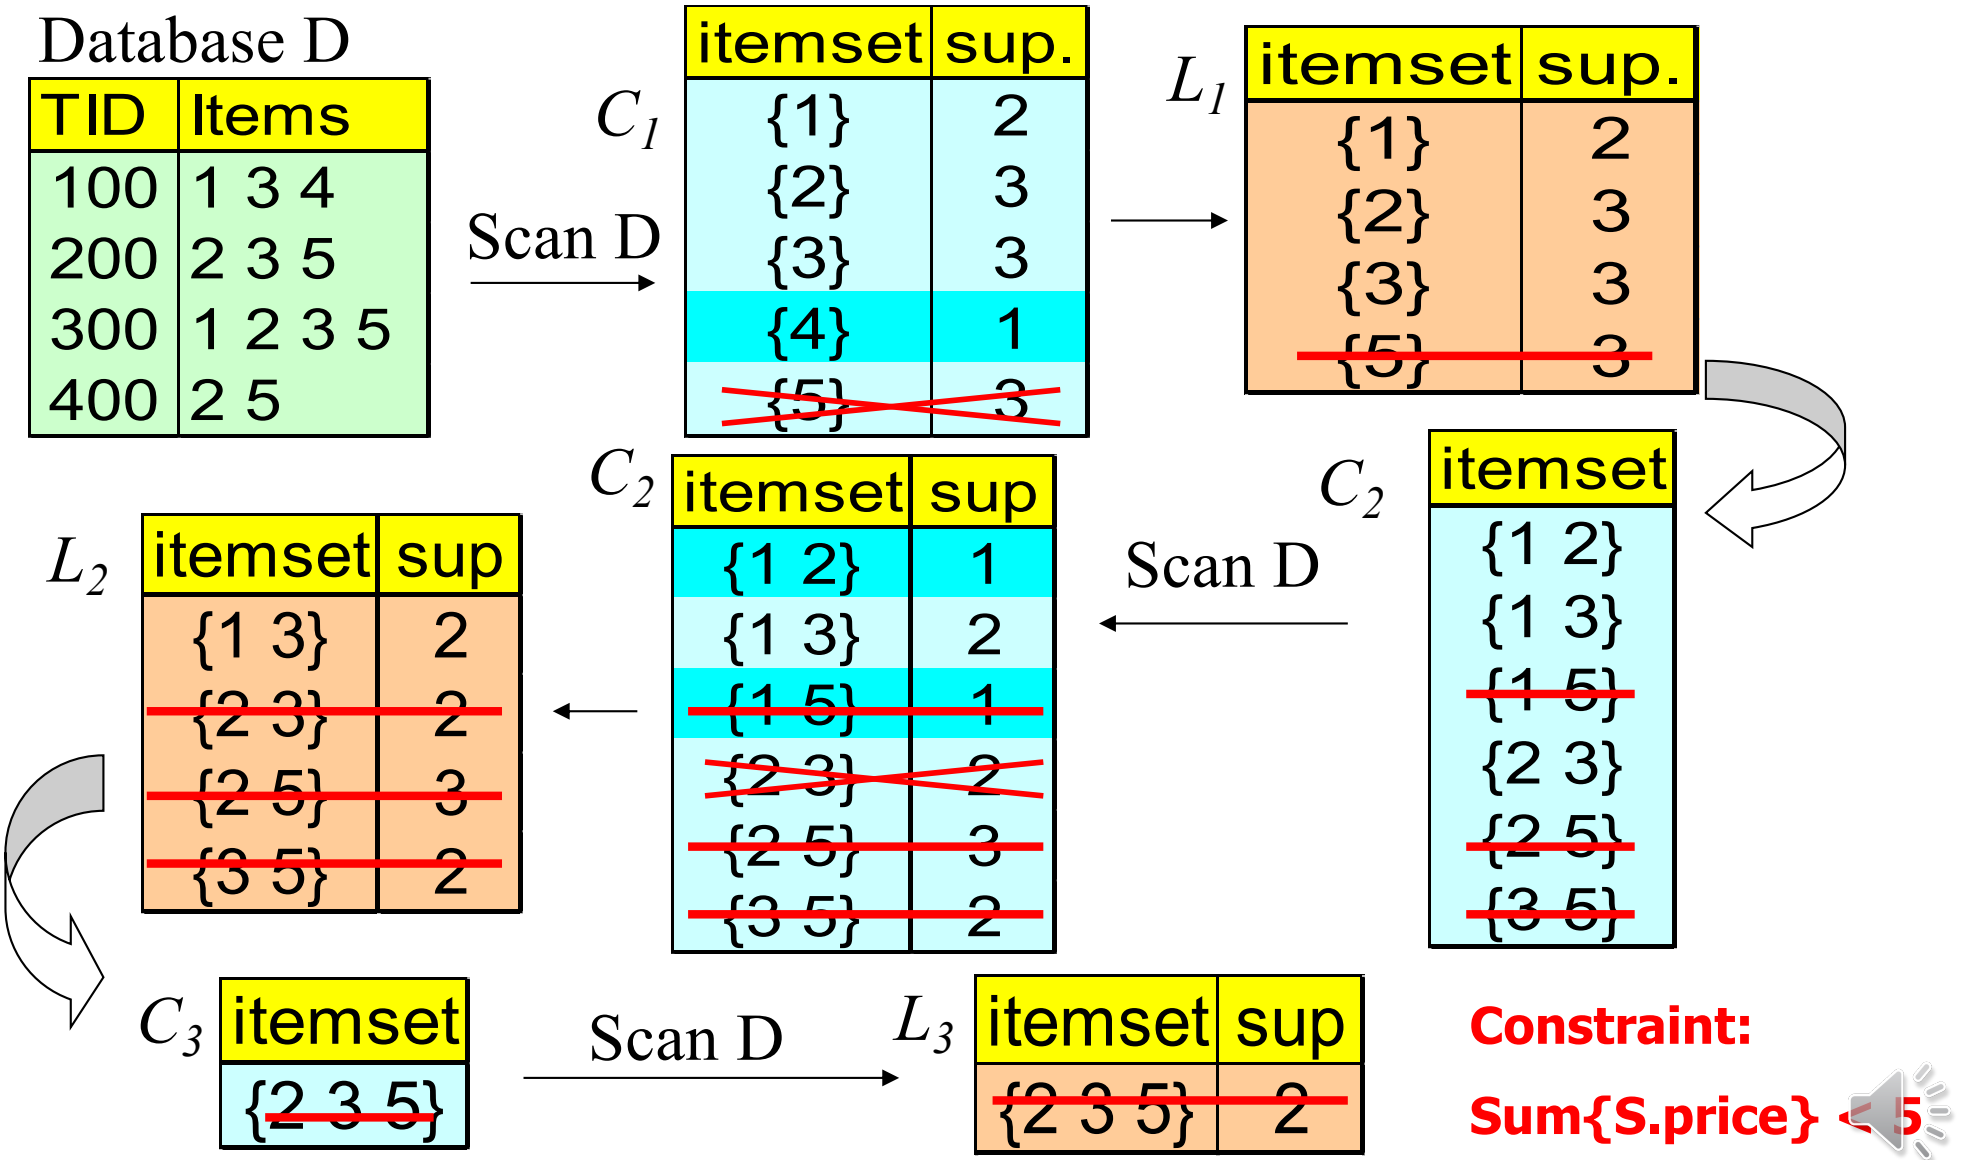

The Constrained Apriori Algorithm: Push a Succinct Constraint Deep

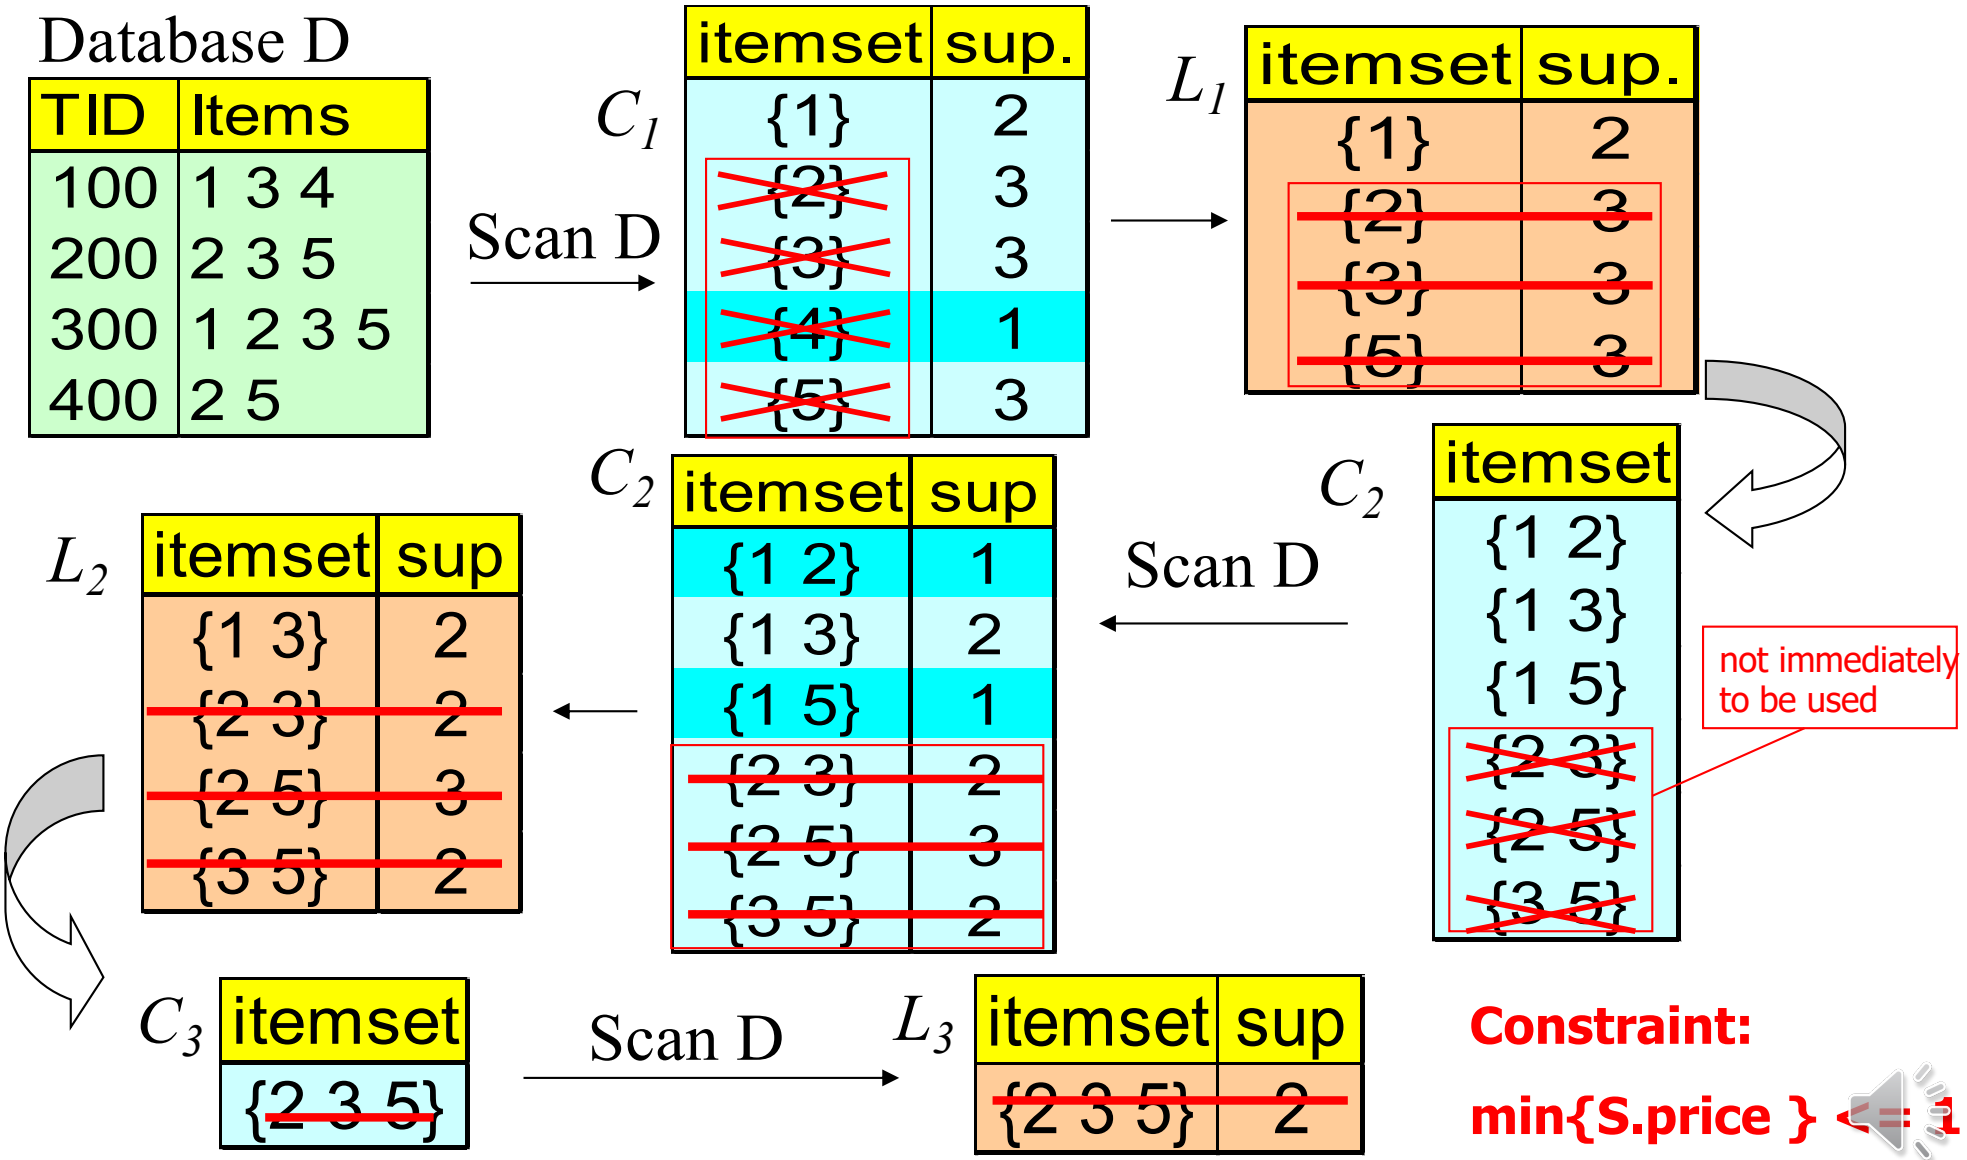

Converting "Tough" Constraints

- Convert tough constraints into anti-monotone or monotone by properly ordering items
- Examine C: avg(*S*.profit) ≥ 25
  - Order items in value-descending order
    - <*a, f, g, d, b, h, c, e*>
  - If an itemset *afb* violates C
    - So does *afbh*
    - It becomes anti-monotone!

TDB (min_sup=2)

TID	Transaction
10	a, b, c, d, f
20	b, c, d, f, g, h
30	a, c, d, e, f
40	c, e, f, g

Item	Profit
a	40
b	0
c	-20
d	10
e	-30
f	30
g	20
h	-10

Chapter 5: Mining Frequent Patterns, Association and Correlations

- Basic concepts and a road map
- Efficient and scalable frequent itemset mining methods
- Mining various kinds of association rules
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

Frequent-Pattern Mining: Summary

- Frequent pattern mining—an important task in data mining
- Scalable frequent pattern mining methods
  - Apriori (Candidate generation & test)
  - Projection-based (FPgrowth, CLOSET+, ...)
  - Vertical format approach (CHARM, ...)
- Mining a variety of rules and interesting patterns
- Constraint-based mining
- Extensions and applications